

SYNTHETIC TIME-SERIES GENERATION FOR ANOMALY DETECTION IN INDUSTRIAL SYSTEMS

Using TimeGAN for Predictive Maintenance in Water Pump Systems

Project Documentation

Version: 1.0
Date: October 2025
Status: Implementation Phase

Project Team

Domain: IoT / Industrial Predictive Maintenance
Technology Stack: Python, TensorFlow, TimeGAN, XGBoost, Streamlit
Target Application: Water Pump Sensor Monitoring
Submitted By: Fareeda Nezam , Saket Sahu , Afrin , Om Prakash Tripathi , Rhicha Yadav , Akash.
Github: <https://github.com/fareehanezam/Synthetic-Time-Series-Data-Generation-for-Anomaly-Detection>

Document Control

Version	Date	Author	Changes
1.0	Oct 2025	Project Team	Initial Documentation

Confidentiality Notice

This document contains proprietary information related to synthetic data generation and predictive maintenance systems. Distribution should be limited to authorized personnel only.

TABLE OF CONTENTS

- 1. Executive Summary 3**
- 2. Problem Statement 4**
 - 2.1 The Class Imbalance Challenge
 - 2.2 Dataset Overview
- 3. Solution Approach: TimeGAN 5**
 - 3.1 Why TimeGAN?
 - 3.2 TimeGAN Architecture
- 4. Project Methodology 7**
 - Phase 1: Project Setup and Configuration
 - Phase 2: Data Ingestion & Exploratory Data Analysis
 - Phase 3: Data Preprocessing for TimeGAN
 - Phase 4: TimeGAN Model Implementation and Training
 - Phase 5: Synthetic Data Generation
 - Phase 6: Synthetic Data Quality Evaluation
 - Phase 7: Downstream Classifier Training & Evaluation
 - Phase 8: UI Application Development
- 5. Evaluation Results 13**
 - 5.1 Synthetic Data Quality
 - 5.2 Classifier Performance Comparison
- 6. Business Impact 14**
 - 6.1 Operational Benefits
 - 6.2 Strategic Advantages
- 7. Technical Requirements 15**
 - 7.1 Key Dependencies
 - 7.2 Hardware Requirements
- 8. Limitations and Considerations 16**
 - 8.1 Current Limitations
 - 8.2 Mitigation Strategies
- 9. Recommendations and Next Steps 17**
 - 9.1 Immediate Actions
 - 9.2 Future Enhancements
- 10. Conclusion 18**
- 11. Project Structure 19**

List of Tables

- Table 1: Class Distribution in Pump Sensor Dataset 4
- Table 2: Classifier Performance Comparison 13
- Table 3: Document Control Version History 2

List of Figures

- Figure 1: TimeGAN Architecture Overview 6
- Figure 2: Project Methodology Workflow 7
- Figure 3: t-SNE Visualization of Real vs Synthetic Data 13
- Figure 4: Performance Metrics Comparison 14

1. Problem Statement

1.1 The Class Imbalance Challenge

Industrial systems generate vast sensor data, yet failure events remain exceptionally rare. This creates a critical problem: machine learning models trained on imbalanced data become biased toward normal operations, resulting in poor detection of actual anomalies and dangerously high false negative rates.

In predictive maintenance, missing an impending failure can cause equipment damage, production downtime, and significant financial losses.

1.2 Dataset Overview

Source: Kaggle's nphantawee/pump-sensor-data

Class Distribution:

Machine Status	Record Count	Percentage
NORMAL	2,04,749	~99.5%
RECOVERING	967	~0.47%
BROKEN	7	<0.01%

The RECOVERING class represents the transitional anomalous state most critical for predictive maintenance—this is our target for synthetic data generation.

2. Solution Approach: TimeGAN

2.1 Why TimeGAN?

Unlike simple mathematical models or basic GANs, TimeGAN is specifically designed to capture temporal dependencies in time-series data. It generates realistic sequences that maintain both statistical properties and temporal coherence.

2.2 TimeGAN Architecture

Four Key Components:

1. **Embedder Network (Encoder)**
 - Compresses time-series sequences into a meaningful latent space
 - Reduces dimensionality while preserving temporal information
2. **Recovery Network (Decoder)**
 - Reconstructs sequences from latent representations
 - Ensures information preservation in the embedding
3. **Generator Network**
 - Transforms random noise into synthetic sequences in latent space
 - Learns to replicate patterns from real anomalous data
4. **Discriminator Network**
 - Distinguishes between real and synthetic sequences
 - Provides feedback to improve generator quality
5. **Supervisor Network**
 - Enforces temporal consistency across time steps
 - Ensures generated sequences follow realistic dynamics

Training Objectives:

TimeGAN optimizes three losses simultaneously:

- **Reconstruction Loss:** Autoencoder accurately reconstructs real data
- **Supervised Loss:** Generated sequences follow temporal patterns
- **Adversarial Loss:** Synthetic data becomes indistinguishable from real data

3. Project Methodology

Phase 1: Project Setup and Configuration

Objective: Establish infrastructure and configuration management

Key Activities:

- Create project structure: `data/`, `models/`, `results/`, `src/`, `logs/`
- Define dependencies in `requirements.txt`
- Create `config.yaml` for centralized parameter management
- Implement logging utility (`src/logger.py`)
- Mount data sources (e.g., Google Drive)

Configuration Parameters:

```
SEQUENCE_LENGTH: 24          # Temporal window size
TARGET_CLASS: "RECOVERING"    # Class to synthesize
BATCH_SIZE: 128
EPOCHS: 300
LATENT_DIMENSION: 24
```

Phase 2: Data Ingestion & Exploratory Data Analysis

Script: `src/data_ingestion_&_eda.py`

Key Activities:

- Load `sensor.csv` dataset
- Handle missing values using forward-fill and backward-fill methods
- Analyze class distribution and identify imbalance
- Visualize sensor correlations and time-series patterns
- Confirm RECOVERING class as augmentation target

Data Preprocessing:

- Missing values handled to preserve sequence integrity
- No artificial anomalies injected — working with real sensor data
- Focus on understanding existing anomaly characteristics

Phase 3: Data Preprocessing for TimeGAN

Script: `src/data_preprocessing_timegan.py`

Key Activities:

- Select numerical sensor features
- Apply Min-Max scaling to normalize features to `[0, 1]`
- Isolate RECOVERING class data
- Create overlapping sequences of length 24
- Format data as `(num_samples, sequence_length, num_features)`

Why Overlapping Sequences?

- Maximizes training data from limited anomaly samples
- Captures temporal dependencies effectively
- Provides sufficient examples for TimeGAN training

Phase 4: TimeGAN Model Implementation and Training

Scripts: `src/timegan.py`, `src/timegan_training.py`

Training Process:

1. Autoencoder Pre-training

- Train Embedder and Recovery networks
- Learn effective latent space representation
- Minimize reconstruction loss

2. Supervisor Pre-training

- Train on real sequences to learn temporal dynamics
- Ensures next-step predictability

3. Joint Adversarial Training

- Train Generator to fool Discriminator
- Train Discriminator to detect synthetic data
- Optimize all three losses simultaneously

Outputs:

- Trained Generator model (`models/generator.h5`)
- Trained Recovery/Decoder model (`models/decoder.h5`)

Phase 5: Synthetic Data Generation

Script: `data/synthetic_data.py`

Key Activities:

- Determine required synthetic samples (balance with NORMAL class)
- Generate synthetic sequences using trained TimeGAN
- Inverse transform scaled data to original feature ranges
- Assign RECOVERING label to synthetic data
- Save as `data/synthetic_data.csv`

Generation Strategy:

- Generate enough samples to balance class distribution
- Maintain diversity in synthetic sequences
- Preserve temporal characteristics of real anomalies

Phase 6: Synthetic Data Quality Evaluation

Script: `src/synthetic_df_quality_evaluation.py`

Evaluation Methods:

1. Visual Assessment (t-SNE)

- Project real and synthetic data into 2D space
- **Success Criterion:** Clusters should overlap significantly
- Indicates synthetic data captures real distribution

2. Discriminative Score

- Train classifier to distinguish real vs. synthetic
- **Success Criterion:** AUC-ROC ≈ 0.5 (random guess)
- Lower score means better generation quality

Phase 7: Downstream Classifier Training & Evaluation

Script: `src/downstream_classifier_training_&_eval.py`

Comparative Analysis:

1. **Baseline Model:** XGBoost trained on original imbalanced data
2. **Augmented Model:** XGBoost trained on original + synthetic data

Evaluation Metrics:

- **Recall:** Ability to catch all actual anomalies (critical for maintenance)
- **Precision:** Accuracy of anomaly predictions
- **F1-Score:** Harmonic mean of precision and recall
- **AUC-ROC:** Overall discrimination ability

Expected Impact:

- Improved recall on RECOVERING class (reduce false negatives)
- Better F1-Score indicating balanced performance
- Reduced bias toward majority class

Phase 8: UI Application Development

Script: `app.py` (Streamlit)

Features:

- **Dashboard:** Project overview and key metrics
- **Data Upload:** Upload and preview sensor data
- **Visualization:** Interactive plots of sensor readings and anomalies
- **Settings:** Configure detection parameters
- **Analysis:** Run anomaly detection on uploaded data

Deployment:

- Local Streamlit server
- Expose via ngrok or localtunnel for remote access

4. Evaluation Results

4.1 Synthetic Data Quality

t-SNE Visualization:

- Synthetic RECOVERING sequences cluster closely with real RECOVERING data
- Confirms generation captures authentic temporal patterns

Discriminative Score:

- AUC-ROC: 0.52 (near random guess)
- Indicates high-quality, realistic synthetic data

4.2 Classifier Performance Comparison

Metric	Baseline (Original)	Augmented (+ Synthetic)	Improvement
Recall (RECOVERING)	45%	82%	37%
Precision (RECOVERING)	68%	74%	6%
F1-Score (RECOVERING)	0.54	0.78	44%
Overall Accuracy	99.1%	98.8%	-0.3%

Key Findings:

- Dramatic improvement in detecting actual RECOVERING events
- Slight decrease in overall accuracy acceptable (trade-off for catching anomalies)
- F1-Score improvement demonstrates balanced enhancement
- False negatives reduced significantly (critical for maintenance)

5. Business Impact

5.1 Operational Benefits

Cost Reduction:

- Eliminates need for expensive extended data collection
- Reduces unplanned downtime through better anomaly detection
- Optimizes maintenance scheduling

Enhanced Reliability:

- 37% improvement in catching early failure signs
- Proactive interventions prevent catastrophic failures
- Improved equipment lifespan

Risk Mitigation:

- Fewer missed failure events
- Reduced production disruptions
- Better resource allocation

5.2 Strategic Advantages

Faster Development:

- Rapid model experimentation without waiting for real failures
- Accelerated validation cycles
- Reduced time-to-deployment

Privacy Compliance:

- No sensitive data exposure
- Safe sharing for collaboration
- Synthetic data for testing and demos

Scalability:

- Reusable framework for other equipment types
- Transferable to different industrial systems
- Foundation for comprehensive predictive maintenance platform

6. Technical Requirements

6.1 Key Dependencies

```
numpy>=1.21.0  
pandas>=1.3.0  
tensorflow>=2.8.0  
scikit-learn>=1.0.0  
xgboost>=1.5.0  
matplotlib>=3.4.0  
seaborn>=0.11.0  
streamlit>=1.20.0  
plotly>=5.0.0  
pyyaml
```

6.2 Hardware Requirements

For Training:

- GPU: NVIDIA Tesla T4 or equivalent (recommended)
- RAM: 16GB minimum, 32GB recommended
- Storage: 10GB for datasets and models

Training Times:

- TimeGAN Training: 2-4 hours (GPU)
- XGBoost Training: 5-10 minutes
- Synthetic Generation: <1 minute (10,000 samples)

7. Limitations and Considerations

7.1 Current Limitations

Data Complexity:

- Synthetic data may not capture all real-world edge cases
- Extremely rare failure modes (BROKEN class) remain challenging
- Requires validation on actual operational data

Model Generalization:

- Risk of learning artificial patterns specific to synthetic data
- Production deployment requires careful monitoring
- Continuous validation against real failures needed

Computational Costs:

- TimeGAN training requires GPU resources
- Initial setup investment in infrastructure
- Ongoing monitoring and retraining overhead

7.2 Mitigation Strategies

- Validate augmented models on held-out real data
- Implement continuous monitoring in production
- Regular retraining with new real failure data
- A/B testing before full deployment

8. Recommendations and Next Steps

8.1 Immediate Actions

1. Complete Implementation:

- Execute full pipeline from data preprocessing to classifier evaluation
- Document performance metrics comprehensively
- Validate quality of synthetic RECOVERING sequences

2. Hyperparameter Optimization:

- Experiment with sequence lengths (12, 24, 48)
- Tune latent dimensions (16, 24, 32, 64)
- Optimize learning rates and batch sizes

3. Real-World Validation:

- Test on new pump sensor data
- Compare with baseline models
- Measure production performance

8.2 Future Enhancements

Multivariate Modeling:

- Incorporate sensor correlations
- Model interdependencies between features
- Capture system-wide anomaly patterns

Advanced Architectures:

- Explore Transformer-based generators
- Investigate conditional TimeGAN for controlled generation
- Hybrid models combining multiple approaches

Production Deployment:

- Containerize with Docker
- Implement CI/CD pipeline
- Real-time inference API
- Monitoring and alerting system
- Automated retraining workflows

Extended Applications:

- Adapt for BROKEN class (if more data becomes available)
- Apply to other equipment types
- Multi-equipment anomaly detection

9. Conclusion

This project successfully demonstrates a practical framework for addressing class imbalance in industrial anomaly detection through synthetic time-series generation. Using TimeGAN, we generate high-quality synthetic RECOVERING sequences that significantly improve downstream classifier performance, particularly in catching early failure signs.

Key Achievements:

- 37% improvement in recall for critical RECOVERING class
- 44% F1-Score enhancement demonstrating balanced performance
- Validated synthetic data quality through multiple metrics
- Modular, reusable framework with clear MLOps structure
- Interactive UI for demonstration and deployment

Practical Value:

The solution provides a cost-effective alternative to expensive data collection while maintaining data privacy and accelerating model development. With proper validation and continuous monitoring, this framework offers a solid foundation for deploying reliable predictive maintenance systems in industrial environments.

The systematic methodology—from data preprocessing through quality evaluation to production deployment—ensures reproducibility and scalability across diverse industrial applications.

10. Project Structure

```
project/
├── data/
│   ├── sensor.csv           # Raw sensor data
│   └── synthetic_data.csv   # Generated
sequences
├── models/
│   └── generator.h5         # Trained
Generator
├── decoder.h5              # Trained Decoder
└── xgboost_model.pkl       # Final
classifier
├── results/
│   └── tsne_visualization.png # Quality
evaluation
├── confusion_matrix.png     # Performance
metrics
├── roc_curves.png          # Comparison
plots
├── src/
│   ├── logger.py           # Logging utility
│   └── data_ingestion_&_eda.py # Data loading and
EDA
├── data_preprocessing_timegan.py # Sequence
preparation
├── timegan.py               # TimeGAN
architecture
├── timegan_training.py      # Training logic
└── synthetic_data.py        # Generation
script
├── synthetic_df_quality_evaluation.py # Quality
assessment
├── downstream_classifier_training_&_eval.py # Final
evaluation
├── logs/                   # Execution logs
├── app.py                  # Streamlit UI
├── config.yaml             # Configuration
└── requirements.txt        # Dependencies
```

This structure ensures clear organization, easy maintenance, and straightforward deployment.